Meridian AI - Data Dictionary

Scope and identifiers

The workload register contains request versions, while the event export contains inference attempts. Neither row population is an account population. Use account identifiers for customer joins, mapping identifiers for rule history, offer identifiers for quotation history, and benchmark batch identifiers for replay observations. No row order carries business authority.

workload.sqlite - request_versions

request_id identifies a logical request. revision orders corrections for that request. account_id can contain case and surrounding whitespace differences. bucket_utc is an hourly September bucket. input_quantity and output_quantity share quantity_unit, which is either tokens or ktokens. outcome is completed or cancelled. cache_eligible is an integer indicator, zero or one. The primary key is request_id with revision. A correction replaces the earlier row; it is not an extra request.

attempt_events.jsonl

Each line is a complete JSON event. event_id identifies an observed attempt and request_id links to logical demand. attempt is the attempt ordinal. Input and output fields contain individual tokens. model identifies the observed historical route. latency_ms is milliseconds. billable identifies provider billing eligibility. Events document troubleshooting activity and are not the authoritative prospective demand register.

customer_forecast.csv

One row represents one account. account_id is the canonical key. route_group links customers to deployment mappings. data_class states standard or restricted; use_case states support, extract or research. oct_factor, nov_factor and dec_factor are absolute volume multipliers, not percentage-point changes. monthly_revenue_usd is contracted revenue for each of the three calendar months.

deployment_options.csv

mapping_id identifies a routing rule and revision orders its retained versions. plan_id is P1, P2 or P3. Each current rule maps one route_group to one model and states its data_class and use_case. one_time_transition_usd repeats the whole plan quote for operational convenience. It is not a per-account, per-group or per-mapping fee. Current rules collectively cover every account group for every plan.

replay_benchmark.csv

batch_id identifies an observation. model and use_case identify its applicable cell. evaluated and passed are counts. p95_ms is a batch latency statistic in milliseconds. retry_attempts_per_request and retry_output_fraction describe incremental retry processing. output_length_factor describes output scaling relative to the observed workload. cache_hit_fraction is the estimated fraction of eligible first-pass input served from cache. Factors are fractions unless the field name explicitly states a count.

infrastructure_offer.csv

Each offer_id covers a pool and prospective month; revision orders quotations. effective_date is an audit date. month is oct, nov or dec. pool and model identify the resource. output_tokens_per_second is measured throughput for one GPU. input_to_output_equivalent converts input processing to output-equivalent work. gpu_hour_usd is the hourly reservation rate. minimum_gpus and maximum_gpus are integer pool bounds. headroom_fraction is permitted utilisation. calendar_hours is the billed calendar period. award_eligible is a true or false field.

Controlling documents

commercial_terms.pdf establishes executed pricing, commitments, eligibility and service requirements. planning_policy.pdf establishes the decision horizon, financial objective and approved calculation conventions. Historical material remains subject to its document status. The request population, contracts and quotation histories should be reconciled before interpreting their combined financial consequences.